



AI/ML Notes 03/08/2021

By Group 5

Various Capabilities / disciplines of an Artificial Intelligence (AI) bot

- 1) Machine learning
- 2) Problem solving (Search, constraint Analysis)
- 3) NLP (Machine Translation, speech Recognition)
- 4) Decision Making (Logic, knowledge, engineering)
- 5) Reasoning (Probabilistic & uncertainty Quantification)
- 6) Robotics (Perception, Action)
- 7) Object Recognition (Vision, Image Processing)

Example

Identify Raga of Music

Computing Info

Sound → Algo → TTS → RAGA

Robotics Part → TTS → Speech engine

What is Artificial Intelligence?

According to John McCarthy it is science and engineering of making intelligent machine and intelligent computer program. It is related to similar task of using machines to understand human intelligence but AI does not have to confine itself to methods that are biologically observable.

Example

Text to speech engine

One can think of that, we can store various words and their corresponding speech signals in database and for translating a sentence, signals of different words can be merged. But the problem here is that different words can have different intonations depending on the context and thus different signal outputs. so storing such large amount of data would require a database that is not physically possible.

So to achieve the goal, machine are trained continuously and certain rules are constructed. Thus the

intelligence is not only restricted to methods that are biologically observable.

Intelligence

Intelligence is a fuzzy term. It is computational part of the ability to achieve goals in the world.

Example

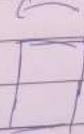
to achieve goals in the world.
example

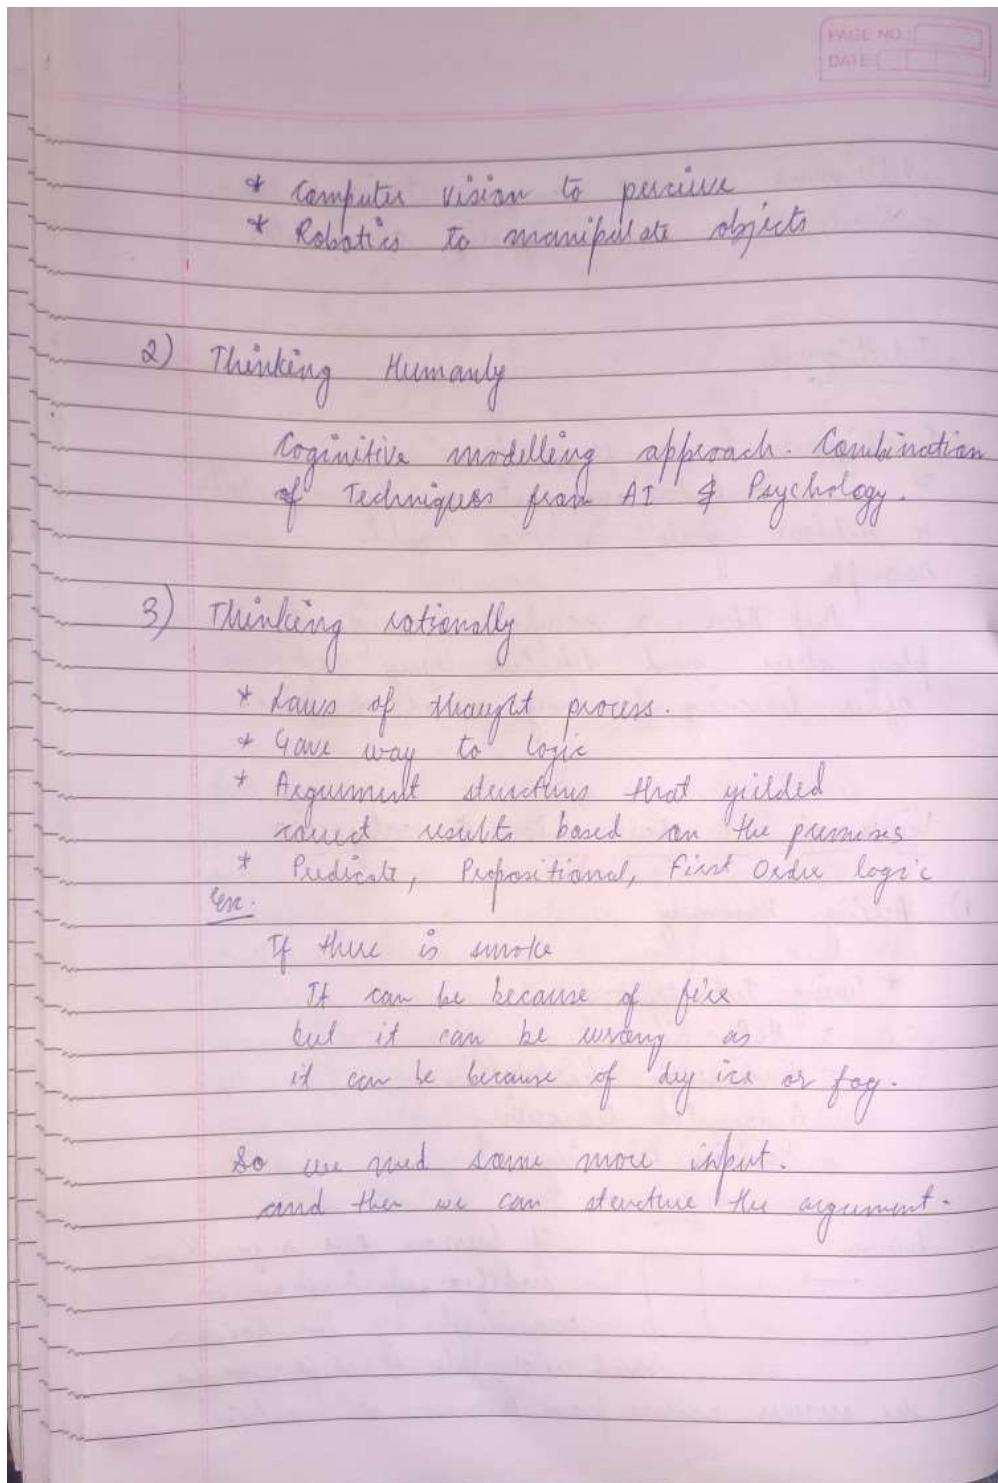
Dub Blue, a computer designed to play chess and defeated Gary Kasparov after learning for long period of time.

Views of AI fall into 4 categories

1) Acting Humanly

- * Turing Test Approach
 - NLP
 - Knowledge representation
 - Automated Reasoning
 - Machine Learning

ans.
human → 
if human pose a question
and if it can human is
convinced by the answer
and not able to diff. into
the answer comes from human or machine.



4) Acting rationally

- * Rational agent approach.
- * To act as to achieve the best outcome
- * And when there is uncertainty the best expected outcome.

* Human-like

- Have to simulate humans intellect and behaviour on by a machine.

• Mathematical Problems (puzzles, games, theorems)

↳ e.g. Sip Blue

• Common sense Reasoning

↳ Don't park if there is no parking sign

• Expert Knowledge: Lawyess, medicine, diagnosis

• social Behaviours

* Rational-like

achieve goals, have performance measure.

↳ for automated cars,
they need to think like rational like

Speech from one lang to another lang ex is missing though raga example is explained, the robotic speech example is missing (how it is stored - why diphones are preferred over words, phonene or syllable), Lombard effect example is missing. Other than the above examples, the notes has all other topics and proper explanation.

-by 106119120 (Sruthi P) and 106119110 (Samiksha Shekhawat)

04th Aug 2021. Class notes prepared by members of group 2 (106119050(H KAILASH) && 106119052(INDRESH P))
Link to pdf: https://drive.google.com/file/d/1E2otsj7JW6FfR_kewthhwVpUj15epxg/view?usp=sharing

4/8/21

Date: / /

Q.No.

Exercise No. /

ASB-31

Subjective Problems : Level / Section -

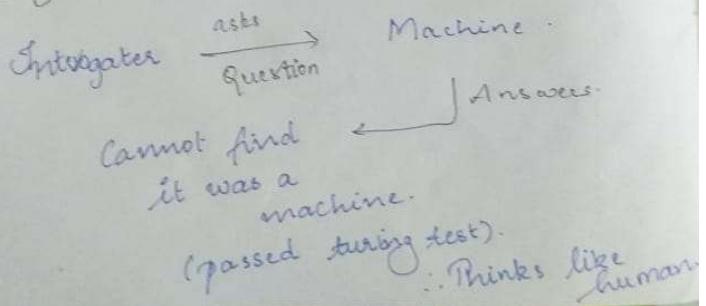
Premise - sequence of argumentative structure.

Computers - Created to do repetitive human work. → evolved to compete with humans. doing human work.

Thinking humanly Acting humanly
Thinking rationally → Acting rationally
 Best outcomes
 uncertainty
 Gurukula,
 Indian philosophy.
Thinking based
on argumentation

Thought process & Behaviour.

Turing test



Requires AI. Knowledge rep., automated reasoning - AI.

Turing → Structured programming

Thinking humanly

Cognitive Science - Psychology - AI.

Thinking rationally

Logic, Problems.

Acting

• Perceive ← get input from env.

• Act

Seq of algo → Output to get acted upon

Perceived

Converted to represented.

AI examples

• Common Sense reasoning.

Smoke means fire.

fire ⇒ Smoke.

Bench mark

& puzzle: Automated parking.

History

- Neural network - 1943.
- Wasn't popular → GPU - 15 yrs back.
(unavailable).
- 1956 - AI name coined.
Logic theorem
principia mathematica theorem
- Summary examples of history.

1956 - Fathers of AI met.

1966-1974. Problem with computation.

1969 - 1979. Knowledge base.
Weak vs Strong.

Recent:- DL / GPU architecture

Bayesian belief network - Uncertainty.

Many more in ppt :)

State of art

Proverb solves crossword.

AI can't replace humans.

Dynamic aspect missing.

Program 1:

Tic-tac-toe

- Data structure approach.

1	2	3
4	5	6
7	8	9

Blank - 0 X - 1 O - 2.

19063 × 9 element vector to work.

39

- Ternary - Decimal.

-Fast in time

Space increasing.

Lot of unnecessary work.

Extend - not possible.

- Data struct approach - 2

Stored as vector.

- Return middle.

- Returns non-corner.

Manipulating only on one vector.

`go()` ← To fetch value. to make move.

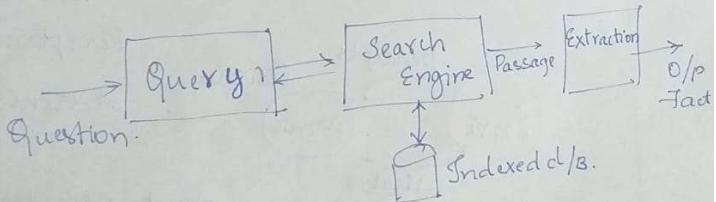
Better than table. (previous).

- State Space tree: $\alpha \rightarrow \beta$.

Doesn't make move which will end losing.

Question Answering.

Question ↘ Definition
Question ↘ factoid.



Eg.: What did Russia do? - How AI approaches.

- Agent & Environments

- Rationality

- PEAS

- Environment types

- Agent types → Search.

Agent

Perceive info. (sensors)



from environment



act upon. (actuators).

Human Agents :- Eyes, Ears . . .

Actuators :- legs, mouth.

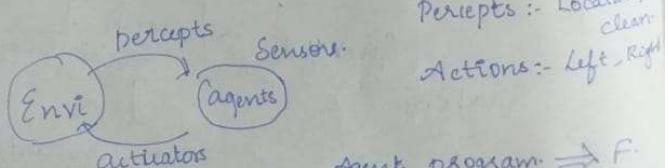
Robotic Agent :- Sensors, motors

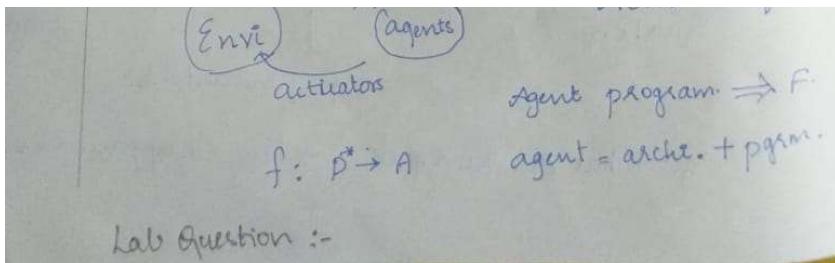
Vaccum cleaner

Percepts :- Location,

clean

Actions :- Left, Right





4th August 2021 – Notes Review by Group 6

Ragavi Vijayaragavan- 10611098
 Shruthi Kumaravel- 106119116

- Acting rationally and the process could have been elaborated further. Information is perceived and converted to a representation. This information is processed by the algorithm and output is delivered in a decipherable form. An example is the AI to find raga.
- In Yale shooting problem, the underlying concept is that a sequence of actions is used to determine if a goal can be reached, with the help of state transition. This wasn't mentioned.
- Expert Systems was not mentioned. Computer systems designed to make decisions like industry experts are called expert systems.
- Otherwise, the notes are concise and cover all the points.

CSPC542 AI / ML
 1) Arunesh Soodish
 106119018
 2) Nitin Benjamin Doshi
 106119088

Difference between AI and ML

- AI is a rational agent with sensors and actuators.
- It can be seen as a block box for ML.
- AI uses sensors to perceive the world and acquire input.
- The ML algorithm uses the given input to decide on an action to affect world.
- The output is realized through the actuators.
- AI maps inputs to output actions through ML algorithms.

Example:
 Automatic robotic vacuum cleaner

:	:
Left	Right

Here the cleaner takes in two inputs, location and status

- If dirty → suck
- Else
 - If right → go left
 - If left → go right

https://drive.google.com/file/d/1xeSoU7qPy_dNah19pMyU8ec4-q1sU97X/view?usp=sharing

PageWork

- 3. Here the robot cleaner senses or perceives two inputs, location and status through its sensors.
- 2. Based on the input, it produces an output using ML algorithms.
- 3. The output is carried out through wheels or the sucking agent, i.e.

sense dirt \rightarrow suck
dust \rightarrow move

Rational agent

- 1. A rational agent is one who picks the right thing even in uncertain situation.
- 2. It picks the right action causing the most success.
- 3. Success is usually calculated using a performance measure.

Example:

The vacuum cleaner can use performance measure based on

- \rightarrow amount of dirt cleaned
- \rightarrow energy consumed
- \rightarrow distance travelled
- \rightarrow Even differentiating between useful and useless sights

Final definition of rational agent:

A rational agent is one that picks the output that maximizes performance, given the evidence provided by the input or percept sequence and the initial knowledge the agent has.

PageWork

Example:

- \rightarrow The robot vacuum will not take diagonal paths usually as it will reduce increase overall distance travelled to cover entire coordinate.
- \rightarrow Higher distance travelled to cover a given area will be lowering performance.

Intelligent agents

- \rightarrow Ability to interact with external world
 - \cdot To receive, understand, act
 - \cdot Eg: speech recognition, understanding systems (Deep Google, Alexa, Amazon echo)
 - \cdot Eg: Image understanding (Computer vision, self driving cars, Biomedical uses)
 - \cdot Eg: Ability to take actions, have an effect.

Example:

An ATM with camera to recognise faces, and a MCC to recognise voice was implemented to help visually impaired people use the ATM.

The cameras also made sure the environment was secure and also identified the user.

If the user was validated and his pin was validated, cash could be withdrawn.

- \rightarrow Knowledge representation, Reasoning and Planning
 \cdot Modelling the external world given input

PogellWork

- Solving new problems, planning, making decisions
 - Ability to deal with unexpected problems, uncertainties

Side notes:

In most ML algorithms around 20% of the data is used to build a model while rest 80% is used to test and the prediction of the model.

→ Learning and adapting

- Continuously able to learn and adapt
- Internal model are constantly updated with each new sequence of percepts.

Implementing agents

- Table look-ups:
 - Primitive algorithm
 - Every possible percept is mapped to an output
 - AI receives a percept, appends its sequence of percepts, and looks up mapped output on the table and acts on it
 - Only possible for limited cases of inputs
 - Example: Vacuum cleaner example.
- Autonomy:
 - All actions are specified,
 - no need in sensing and in autonomy
 - self sufficient

PogellWork

- Omnicience
 - Knows the outcome of an action carried out
 - Example: Traffic signal
 - If red → stop
 - If green
 - If previous person was
 - move
 - else wait

We need AI to have autonomy, omniscience, and learning.

- Learning:
 - Infers information from percepts

Example: Vacuum cleaner might use camera to learn coordinates of a room.
If a new furniture is added it will adapt and adjust the coordinates to learn.
It decides on actions based on coordinates and percepts.

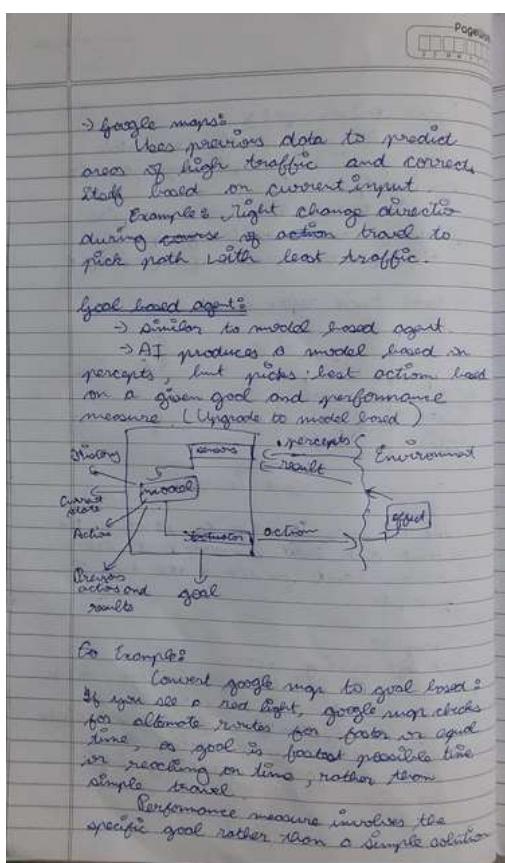
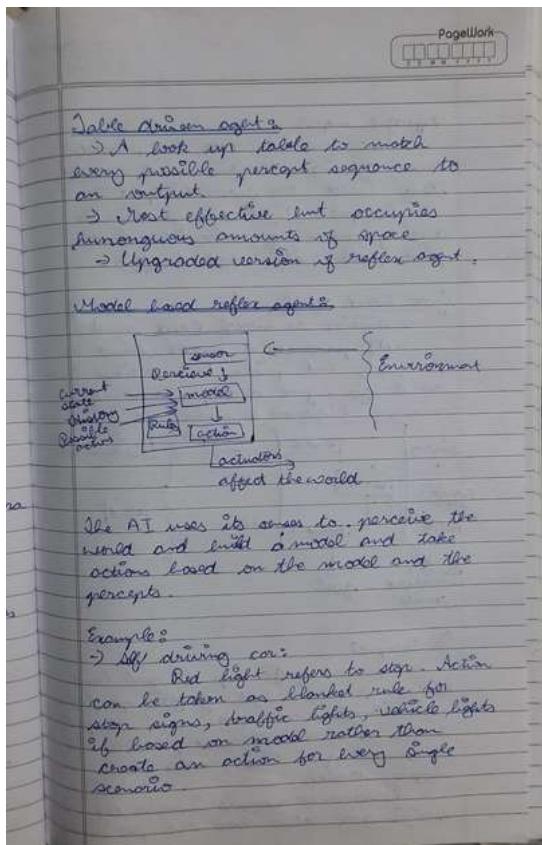
Reflex agent (Simple)

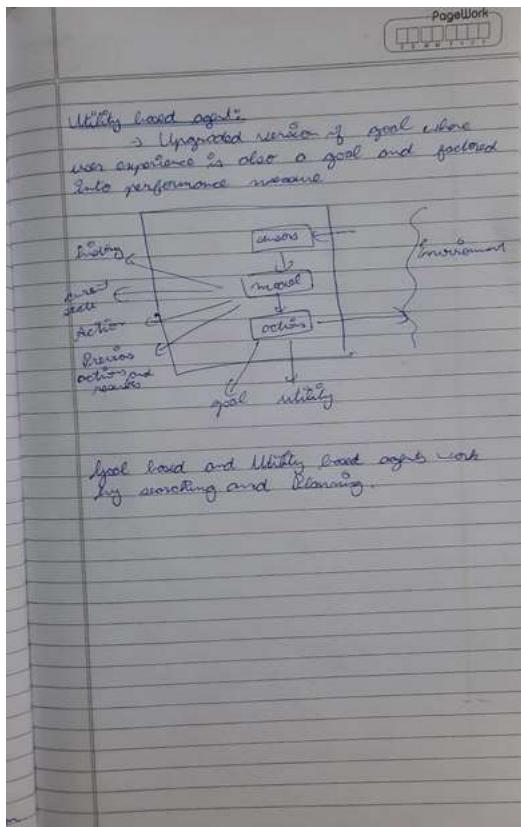
An AI that works solely based on current input is a reflex agent based AI:

Eg: Vacuum cleaner

```

    If dirty → suck
    else if left → go right
    if right → go left
  
```





Group 13 evaluation 9th August

Anirudh VS - 106119012

Rinish Sam I - 106119102

- Explanation is missing for agents with autonomy: The agent should be able to design and act, given a new environment.
- Example for agents with autonomy w.r.t the vacuum cleaner scenario is missing.
- For model training 80% of data is used to train the model and 20% is used to test the model. These values are mentioned in correctly.
- Example for utility based model is missing.
- Apart from this all the points are covered.

Class notes of 10th August 2021. Prepared by Group 15: Madhavan B(106119068) and Sakthi Prasath S(106119108).

Link to PDF:

[https://drive.google.com/file/d/1Mi1boRfKvp7Z05ULsJit8pT0aH375X-J/view?
usp=sharing](https://drive.google.com/file/d/1Mi1boRfKvp7Z05ULsJit8pT0aH375X-J/view?usp=sharing)

* Goal, Utility based Agents require lot of searching and planning.

, Fully observable environment vs Partially

↓
Complete environment is ↓
available partial env.

The other one is ~~not~~ observable env.

→ Google Maps: Fully observable env.

↳ If senses: complete env. at any particular point of time.

→ Autonomous driving vehicle: Fully observable env.

→ Vacuum cleaner: Partially observable env.

* Autonomous vehicle is a dynamic agent.

it must be fully observe the env. to avoid risk.

* If autonomous vehicle gets ~~sense~~ many may not sense the full env. So it is partial at some times.

→ Single agent vs Multi agent env.

↳ Crossword puzzle solver.
(All the hints and other data are present with itself)

↳ Autonomous vehicle
(Different agents give input to an agent)
Chess

→ Deterministic vs stochastic

↓
If the next state depends upon the current state and the action.

↓
If the next state doesn't depend upon the current state and the action.

Stochastic can be thought as partially observable.

Deterministic can be thought as fully observable.

be called

* Taxi driving can be ~~seen~~ as stochastic environment.

* Episodic vs Sequential

↓
Action is divided into atomic components

* next action doesn't depend on previous action. (Ex: gaming, vacuum cleaner)

↳ small components integrated into a big one.
(Ex: Garden sprinklers).

* Static vs Dynamic

* Ex: crossword puzzle, garden sprinkler.

* Ex: games.

- * Discrete vs Continuous
Let us consider the environments have states and time is a factor.
 - * If env is defined at some points in time
 - * Ex: chess, games.

- * Known vs Unknown
Agent knows
- * Outcomes are available for all actions.
- * Ex: Tic Tac Toe game with all possibilities recorded.

To evaluate ~~an agent~~, we have a measuring scale known as PEAS.

P - Performance measure
E - Environment S - Sensors.

	Environment Types			Partly
Solitaire	Yes	Yes	Internet shopping	Taxi
Observable	Yes	No	No	No
Deterministic	Yes	No	Partially	No
Episodic	No	No	No	No
Static	Yes	Semi	Semi	No
Discrete	Yes	Yes	Yes	No
Single Agent	Yes	No	No	No
	(except auction)			Tutor
Observable	Partial	Interactive	No	
Deterministic	No	Stochastic		
Episodic	No	sequential		
static	Dynamic	Dynamic		
Discrete	Continuous	Discrete		
Single Agent	Single	Multiagent		

- * The real world is (of course) partially observable, stochastic, sequential, dynamic, continuous, multi-agent.

PEAS for English Tutor

Performance Measure: Student's performance ^{Oral} written
Env : set of students, Testing agency

A: Marks,

S: Keyboard Entry / Oral English capture

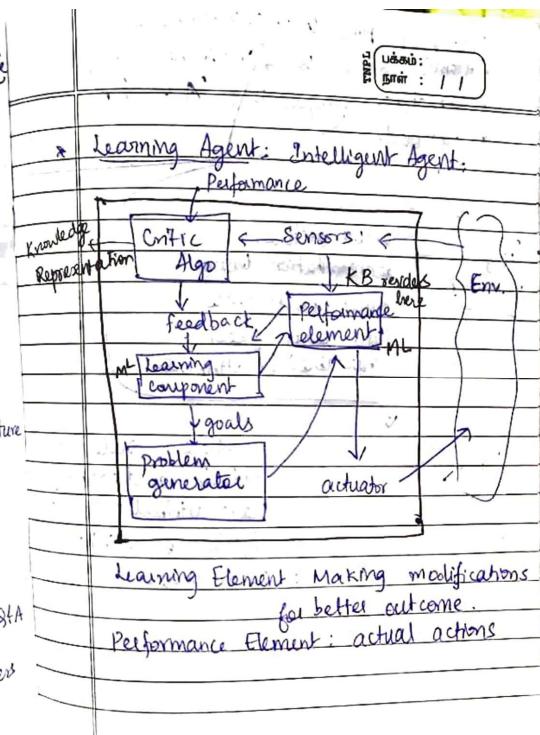
PEAS for Medical Diagnostics

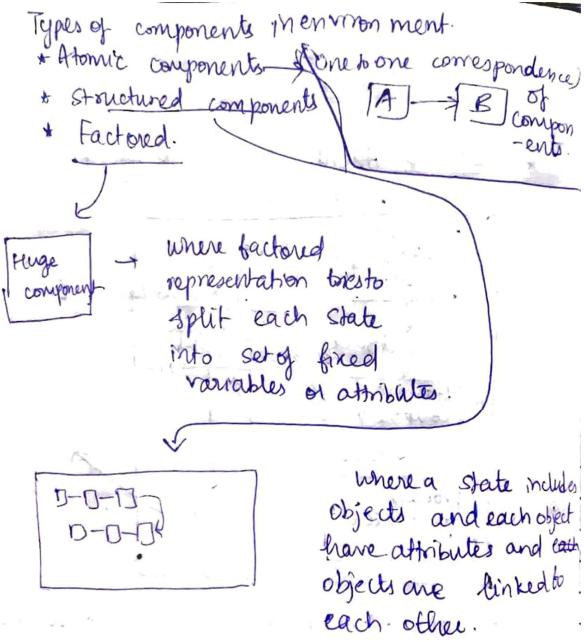
P: Healthy patient, low cost

E: Hospital, Patient, Staff, Lab.

A: Diagnostic report, Treatment, Refund, Q/A

S: Entry of symptoms, Patient's answers





Example for atomic representation: Game playing |

Hidden markov models. This is based upon HMM

Markov Decision Process (MDP) [will learn about it later.]

Example for factored representation:

- Constraint Satisfaction Problems
- Problems based on propositional logic (CSP)
- Bayesian networks
- ML Algorithms

Examples for structured representation:

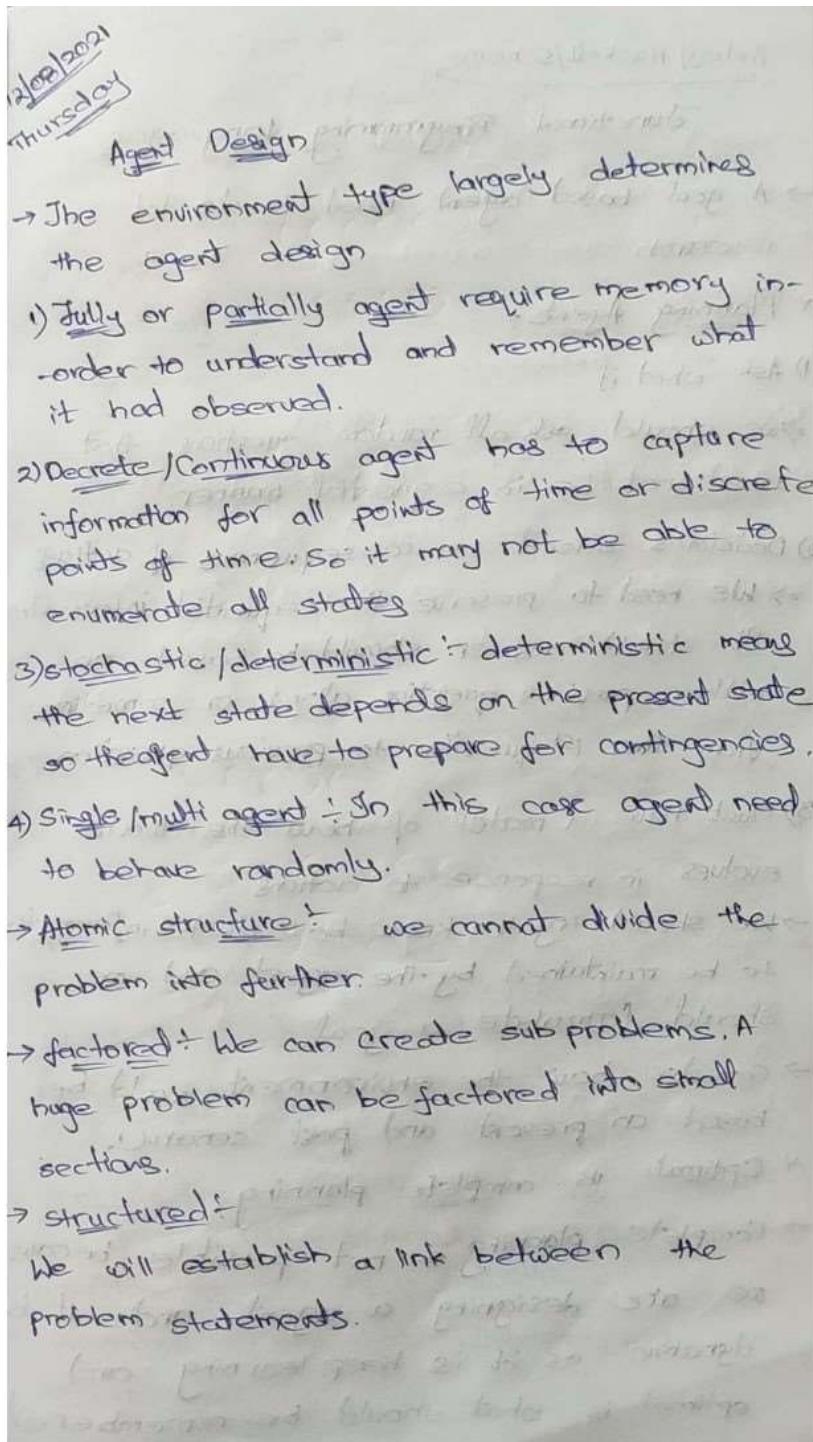
- First order logic
- Natural language Understanding (NLU)
- Knowledge based learning (KBL).

In structured representation knowledge representation plays a major role. We would work with 'ontology'.

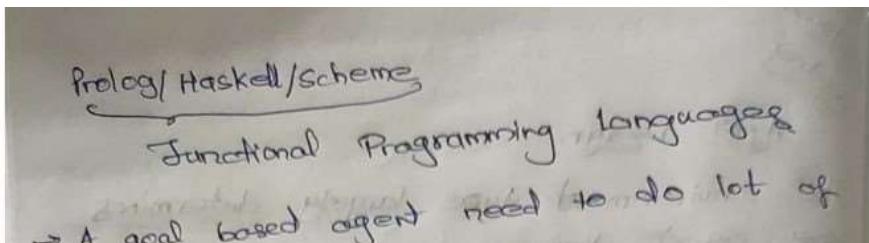
Ontology is the connection between entities and relationships.

- Atomic components - (explanation)
We cannot divide the environment further(missing)
- Apart from this all the points are covered

12/08/2021 Notes – Group 20 (106119136,106119132)



Scanned with CamScanner



→ A goal based agent need to do lot of research

* Planning Agents!

1) Ask "what if"

→ we should ask all random questions that need not be in sequential manner.

2) Decision's based on consequences of actions

→ We need to preserve the sequential information. It should answer should be based on either previous question alone or sometime from the 1st question to previous question.

3) Must have a model of how the world evolves in response to actions.

→ We should try to know how much information to be maintained by the agent and we should formulate a goal.

→ Consider how the environment would be based on present and past scenario.

* Optimal vs. complete planning

→ Complete planning is not possible, because we are designing a agent and it is dynamic as it is keep learning and optimal is what should be remembered

Scanned with CamScanner

* Goal test!

→ Test to verify the given state is final state or not

* Path cost!

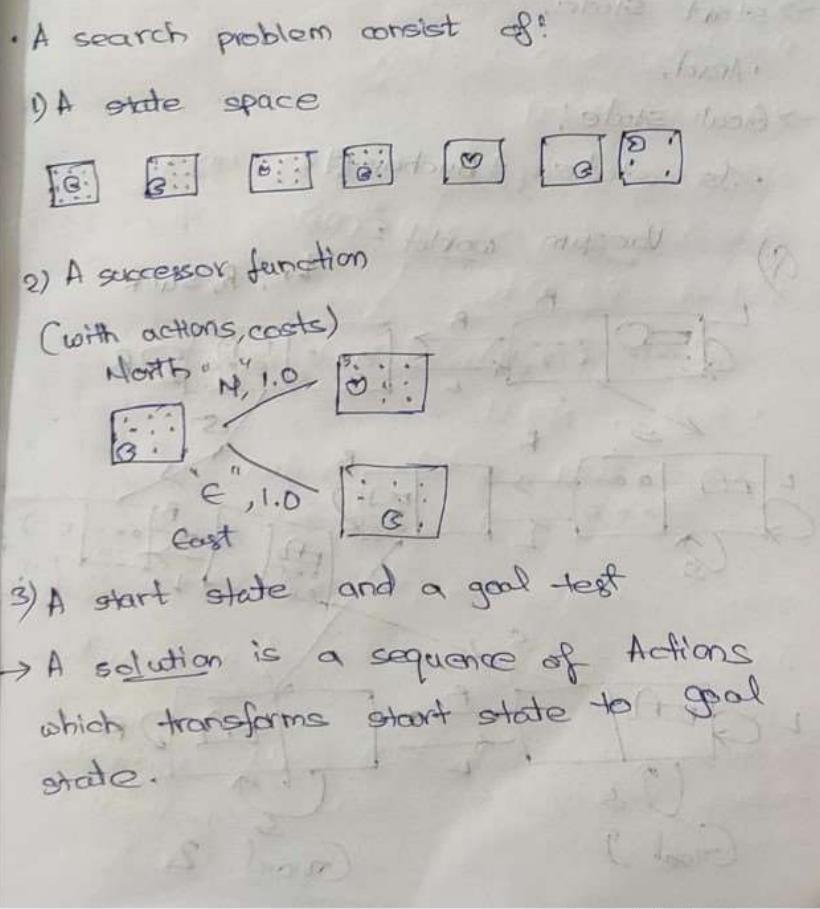
→ Function to assign a numeric cost to each path

* Solution!

→ The one that leads to the final state from the goal state.

Search Problems

• A search problem consist of:



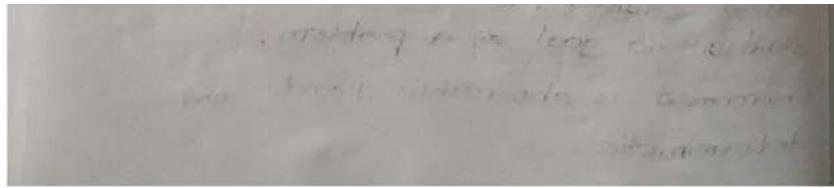
Scanned with CamScanner

- and what should not be
- * planning vs. replanning
 - Plan ahead and replanning is working on the already planned information to do replan.
 - planning agent and replanning agent works together.
 - e.g. pac man
 - In replanning we reduce the time by making it not to move in the empty space much.
 - Search problems are used by Problem solving agents.
 - * Problem Solving Agents
 - Uses atomic representations: actors are considered as whole or no

- Uses atomic representations: states are considered as whole or no internal structure visible to the problem. no possibility to split them further.
- Goal formulation: Based on current situation and agent's performance measure, it is done.
- Problem formulation: We have goal and current situation. We need to convert current situation to goal as a problem.
- Environment is observable, discrete and deterministic

Scanned with CamScanner

- Process of looking for a sequence of actions is called search.
- Search take a problem as input and returns a solution in the form of action.
- Actions are carried out in the execute-phase.
- Search Problems:
 - 1) Initial state - starting point
 - 2) Actions - Description of the possible actions give the current state.
 - 3) Applicable action - For a particular state some of the action will be done.
 - 4) Transition Model - A description of what each action does.
 - 5) Successor - state reachable from a given state.
 - 6) State space - set of all states reachable from initial state by any sequence of actions.
 - 7) Graph - Nodes are states and the links between nodes are actions.
 - 8) Path - Sequence of states.



Scanned with CamScanner

* Goal test:
→ Test to verify the given state is final state or not

* Path cost:
→ Function to assign a numeric cost to each path

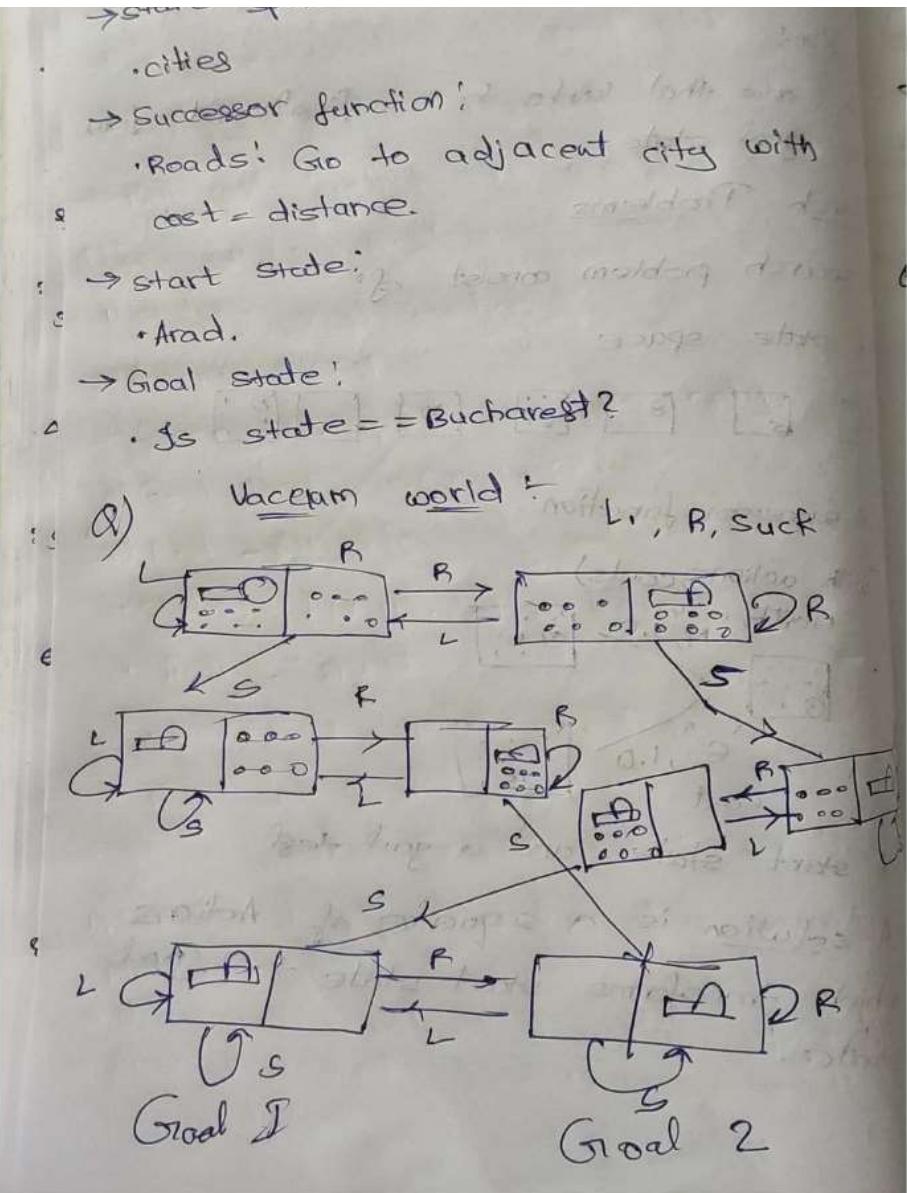
* Solution:
→ The one that leads to the final state from the goal state.

Search Problems

- A search problem consists of:
 - 1) A state space
 - 2) A successor function
(with actions, costs)
North "N", 1.0
East "E", 1.0
 - 3) A start state and a goal test
- A solution is a sequence of Actions which transforms start state to goal state.

Scanned with CamScanner

Eg: A map consisting of cities as nodes and edges as cost for travelling between them.
→ state space:
• cities



Scanned with CamScanner

states: 2×2^2 Agent $\rightarrow 2$

$\Rightarrow 2 \times 2^2$ possible states

$n \rightarrow n \times 2^n$

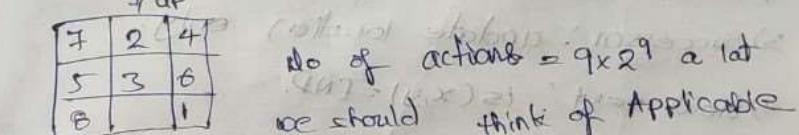
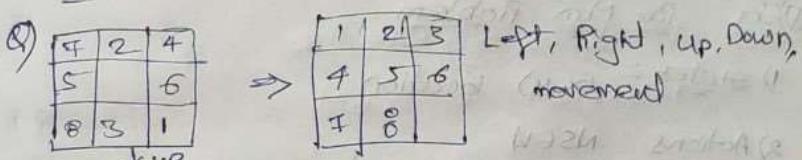
Initial State:

Action: L, R, S

Transition Model:

Goal test: all the squares are clean

Path test: each step cost 1 \Rightarrow no of steps



7	5	6
3		
8	1	

No of actions = 9×2^7 a lot
 we should think of Applicable
 actions.

↓ left

7	2	4
5	3	6
B	1	

Q) 8-queens problem

Initial state? No queen

action! Add a queen such that No queen
 attack each other

transition state!

Scanned with CamScanner

Q) Route Finding:

state: location.

→ initial state: user?

→ Actions: any mode of transport.

→ transition Model: state result on go.

→ from one state to next state

→ Goal test: Have I reached destination?

→ Path cost: Petrol (or) Calories

Q) Pac Man Problem

1) states: (x, y) location

2) Actions: NSEW

3) Successor: update location only.

4) Goal test: $(x, y) = \text{END}$.

Q) Eat - All - Dots.

• states: $\{(x, y), \text{dot boolean}\}$

• Actions: NSEW

• successor: update location, and
possibly a dot boolean

→ state space Graphs and search Trees

1) Mathematical representation of a search problem

→ Nodes are (abstracted) world configurations

→ Arcs represent successors (action results)

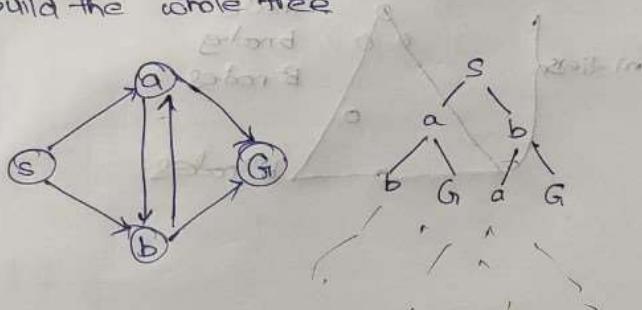
→ The goal test is a set of goal nodes

Scanned with CamScanner

- In state space graph each state occurs only once.
- We can rarely build this full graph in memory (it's too big), but it's a useful idea.

Search Trees:

- A "what if" tree of plans and their outcome
- The start state is root node.
- Children correspond to successors.
- Nodes show states, but correspond to PLANS that achieve those states
- For most problems, we can never actually build the whole tree



* searching with a search tree

- Expand out potential plans (tree nodes)
- Maintain a fringe of partial plans under consideration.
- Try to expand as few tree nodes as possible

Scanned with CamScanner

Depth-First Search

strategy: expand a deepest node first

Implementation:

Fringe is a LIFO stack

Properties of Search Algorithm

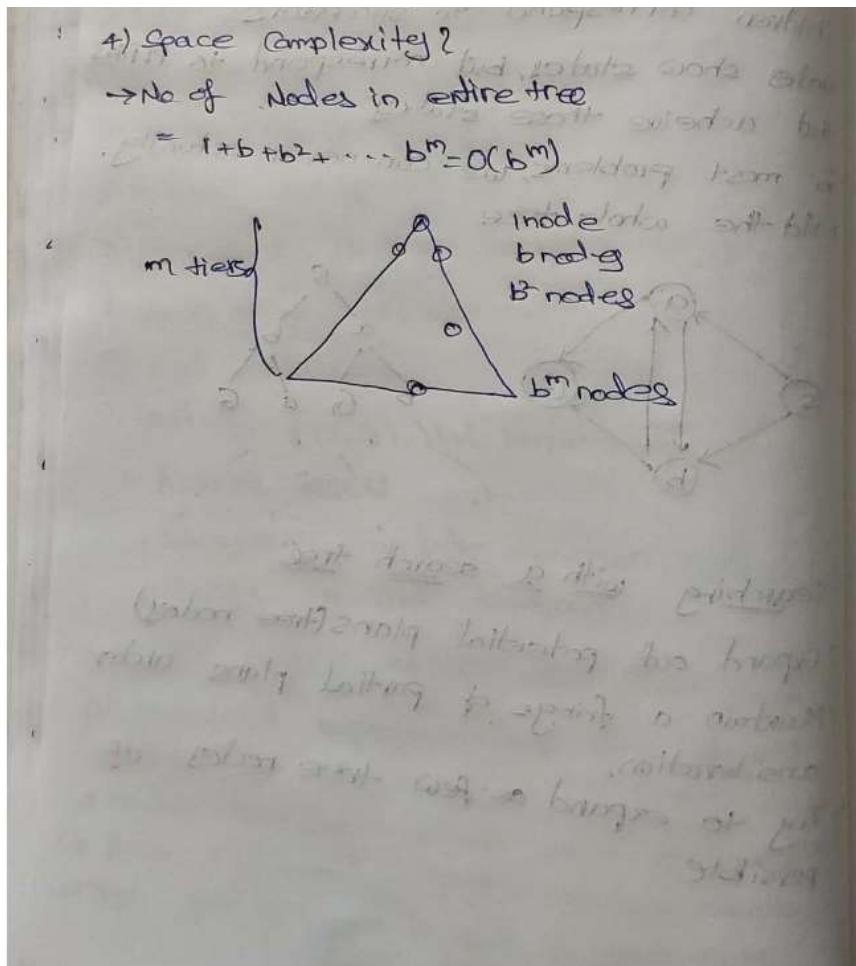
1) Complete:

2) Optimal:

3) Time Complexity?

4) Space Complexity?

→ No of nodes in entire tree.



Scanned with CamScanner

12th August 2021 – Notes Review by Group 18

Pranav Somaiah- 106119094
Vaasu Gambhir- 106119140

- Prolog was classified under functional programming languages while in reality it is a **logic** programming language, also the example of Lisp was missed out on.
- Under Planning agents, the chatbot example wasn't mentioned for asking "What if" and under Problem Solving Agents, the map example wasn't mentioned(Google Map to get directions).
- Lex and Yacc, and how a problem is looked upon in a rule-based programming language wasn't mentioned(like how we state the problem instead of approaching it sequentially).
- Fringe is a LIFO stack that contains vertices of the path that is currently being visited, this was not specified fully.
- Important ideas of the DFS search were not given fully (they are: having fringe, an exploration strategy and expansion). That along with steps of finding an optimal solution using DFS weren't given.
 - > Initialise starting state with init state,
 - > If there is no way to expand and/or goal state is not reached, return failure, else return success
 - > Else keep exploring candidates which can be further expanded.
- It isn't specified that we want to complete the problem that consumes least time and space and the solution should be optimal(under DFS)
- Apart from this the notes are well written and cover major portions of lecture and PPT.



AIMLnotes1
6_08

16th Aug 2021. Class notes by group 23 (106119042(G.Subhasree) &&
106119130(Tanu Gangrade))

Link to pdf: [https://drive.google.com/file/d/1qIE_K259BX4E1nWwhKOf3c7_DHjeCtkT/view?
usp=sharing](https://drive.google.com/file/d/1qIE_K259BX4E1nWwhKOf3c7_DHjeCtkT/view?usp=sharing)

16 | 8 | 21

Recap : Properties of Search Algorithm:

- Complete
- Optimal
- Time Complexity \rightarrow Should be minimal
- Space Complexity

General Tree Search:

Fringe : Vertices which are part of tree, but not visited yet.
 \rightarrow given

function TREE-SEARCH (problem, strategy) returns a solution, or failure

Initialize the search tree using the initial state of problem

loop do \rightarrow no soln.

if there are no candidates for expansion return failure

choose a leaf node for expansion according to strategy.

if node contains goal state, return corresponding soln.

else expand node & add resulting node to search tree

end

Imp ideas:

Fringe

Expansion

Explore" strategy.

Main question: Which fringe nodes to explore??

e.g. Tree Search [Bfr. last class]

\hookrightarrow search strategy depending on how we search

Uninformed Search Strategy [Technique]:

- Goal state not known [ahead of time]
- Keep expanding, till the goal \Rightarrow done

\hookrightarrow DFS

\hookrightarrow BFS

\hookrightarrow Uniform Cost Search

\hookrightarrow Depth Limited Search

\hookrightarrow Iterative Deepening Search

\hookrightarrow Bidirectional Search

To be evaluated
for:

- Completeness
- Optimality
- TC
- Space Complexity

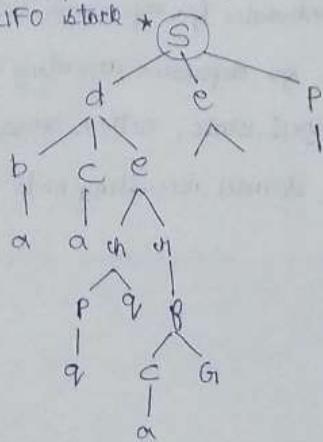
DFS:

last class eg.*

about Gram 5

DFS:

- last class eg.*
- expand deepest node first
- Fringe: LIFO stack *



- Start from S
- [d, e, p are fringe vertices]

- Pick b, go till a
↓
not goal
↙
backtrack
- continue as on
⋮⋮⋮⋮
- reach G₁: Goal state

Soln.

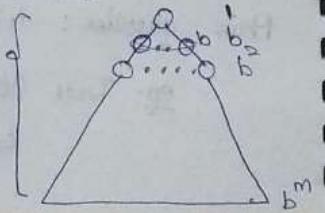
Go as deeper as want

Eval: Soln. in least cost: not guaranteed
[Complete \Rightarrow Optimal]

discussed last class

{ TC: order b^m
SC:

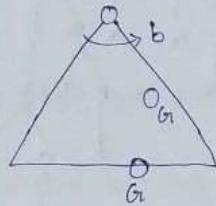
m tiers



- No. of nodes: $1+b+b^2+\dots+b^m = O(b^m)$
 - \downarrow upper bound
 - b: branching fac.
 - m: max. depth
- fringe vertices : in stack
 - ↳ grows big
 - [including repeating nodes]

→ Which nodes are expanded?

- Some left prefix of tree
- whole tree could be processed
- m: finite $\Rightarrow O(b^m)$ time



→ Space of fringe?

- Only shows siblings on path to root $\Rightarrow O(bm)$

→ Is it complete? [get soln. for sure] DFS directed graph

- Not always
 - \downarrow sometimes can't reach
- m could be ∞ , \Rightarrow prevent cycles

→ Is it optimal?

- Not complete \Rightarrow Not opti
- finds leftmost, regardless of cost

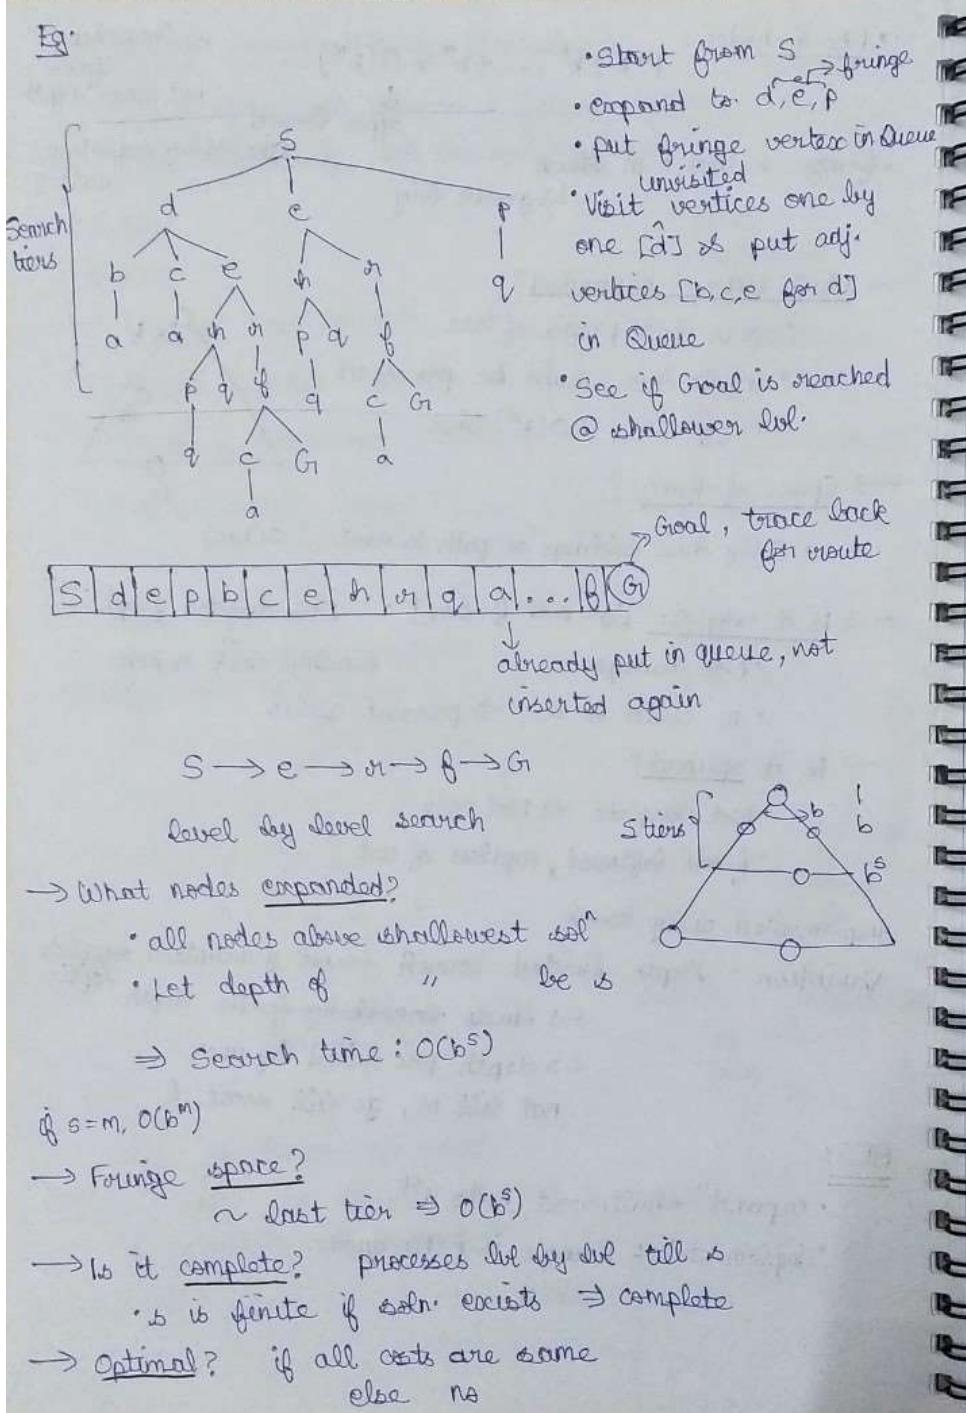
Implemented using stack

Variation: Depth Limited Search \rightarrow not guaranteed complete / opti.

- ↳ limits branching factor depth
- ↳ depth pre fixed by user
- not till m, go till some l

BFS:

- expand^a shallowest node 1st
- Implement^a : fringe is FIFO queue



Modification of BFS: UCS

DFS vs BFS:

BFS outperform DFS:

- less connected graph, goal @ shallower depth

DFS outperform BFS:

- if graph is completely connected, get soln.
- in left half itself in DFS
- [every node connected to every other]

Iterative Deepening Search:

- Grid strategy
- Combines BFS & DFS
- Idea: get DFS's adv of space with BFS's time / shallow adv.
- Deepen only after exploring shallowness
- Not a strategy of BFS, not stack strategy of DFS

Increases limit of algs : from 0, 1, ...

uses depth limit concept for limit of do DFS

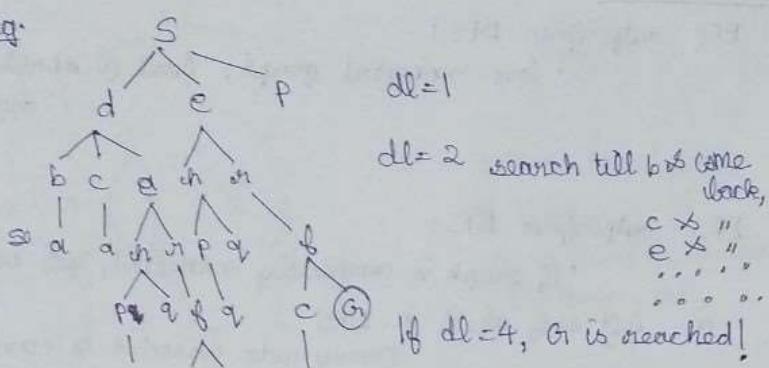
- Run DFS with depth limit 1. If no soln: ↗ [nothing but BFS]
 → " " " " " 2. " " "
 [all vertices not put in stack,
 dl=2, only till 2 its put]
 → " " " 3. ... , .

* Redundant wastefully? not costly, complete & opt-sln.

as it combines BFS & DFS

• most work in deepest sol search ⇒ not so bad!

e.g.



• limiting depth: BFS incorporated

• called iterative as dl is not pre fixed

- limiting depth: BFS incorporated
- called iterative as dl is not pre fixed

Cost Sensitive Search: [Uninformed]

- BFS opt. only when costs or same
finds shortest path in terms of no. of actions

↳ To find least cost path

Uniform Cost Search

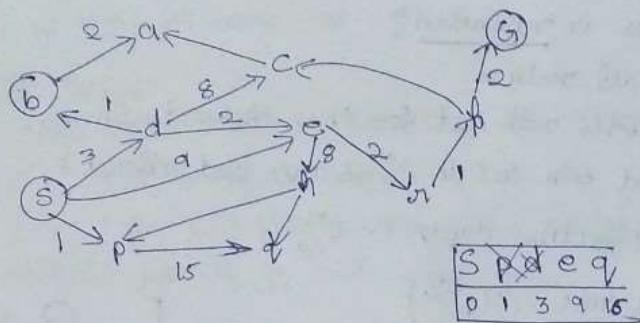
Uniform Cost Search [UCS]:

- Variation of BFS
- gets least possible cost
- Not queue strgy; greedy expansion of fringe vertices to choose path
- Put vertices in queue, based on cost

Strgy: expand cheapest node 1st

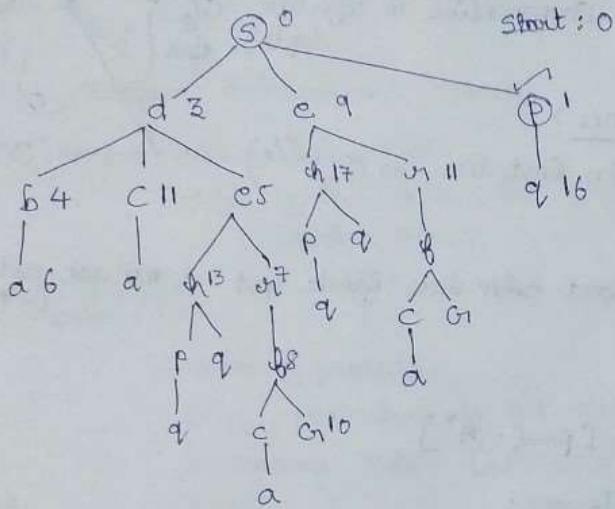
Fringe is priority [cumulative cost] queue.

Eg:



start: 0

Go with lowest cost
BFS with first node
visited being lowest cost
node cumulative



S	X	X	e	a
0	1	3	9	16

visit

d

S	X	X	a	c
4	X	5	6	11

A	B	C	D
8	10	11	13

A	B	C	D
6	7	11	13

G	9	C	In	g
8	10	11	13	16

↓

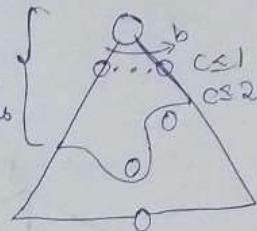
reached Goal with min cost

~ to ~~forwards~~ Dykstra's alg to find
[only shortest path
edge cost] from S - Gr

↳ What nodes u expanded?

- not all vertices
- all nodes with cost less than cheapest soln.
- If that soln. cost is C^* & arc cost atleast ϵ ,
⇒ effective depth $\sim C^*/\epsilon$
- Takes time $O(b^{C^*/\epsilon})$

[exponential in effective
depth]



↳ fringe space?

- roughly last tier, $\Rightarrow O(b^{C^*/\epsilon})$

↳ complete?

- if best soln. has finite cost & min arc cost is +ve,
yes

↳ Optimal?

yes [proof: A*]

Uniform Cost Issues:

- explores increasing cost contours
- complete & opti.
- cheapest cost wise when compared to other algs

Issues: uniformed & ...

- explores in all dir'
- No info abt goal loc'

to be covered with informed strat.

- Demas:
- 1: UCS with some knowledge about goal
 - 2: UCS
 - 3: BFS as the path goes via first deepest node

- BFS is UCS with same cost

The One Queue:

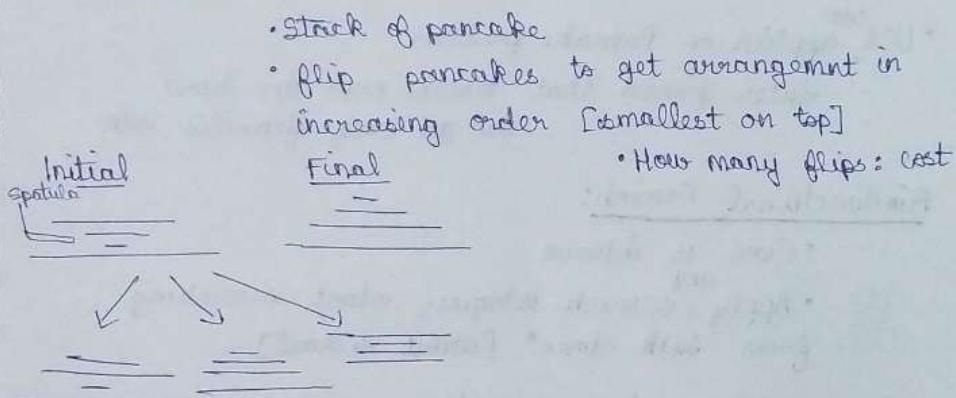
- All algos u same except fringe strates:
 - all fringes or priority queues
 - for BFS & BFS, can avoid $deg(n)$ overhead from

- all fringes or priority queues
- for DFS vs BFS, can avoid $\deg(n)$ overhead from actual priority q_i , with stacks $\propto q$
- Can code an implementation with variable queuing obj.

*One more strgy.: Bidirectional search

eg. Google map: ocean b/w 2 pts.
says go thr. ocean b4 6 yrs.
[walk in ocean]

Pancake Problem:
[Chapathi / deer]



- In 1978, Gates & Papadimitriou did:
for permut. σ of $1 \dots n$, $f(\sigma)$: smallest no. of prefix reversal
for $\sigma \rightarrow$ identity permu.
- $f(n)$: largest $f(\sigma)$ for all σ in (symm. grp.) S_n
we show, $f(n) \leq (5n+5)/3 \Rightarrow f(n) \geq 17n/16$
 n : multiple of 16

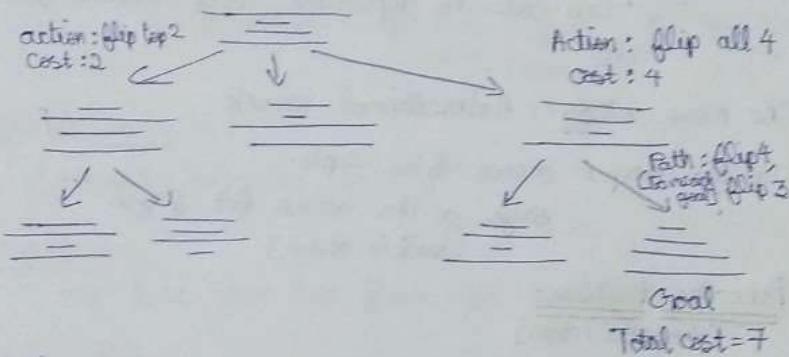
If each int. is required to participate in even no. of reversed prefixes, corresponding $g(n)$ obeys:

$$3n/2 - 1 \leq g(n) \leq 2n + 3$$

III b/c to:

$2^n - 1$ moves for Tower of Hanoi

→ State space tree is based on gen. tree search



* UCS^{was} applied on Pancake problem

Total cost = ?

* UCS was applied on Pancake problem
 later proved that Pancake prob was hard
 can get only feasible soln.

Bi-directional Search:

- Goal is known
 - Apply ^{any} search strategy, start searching from "both dirns" [start & goal]
 - Soln., when u meet
 - Time is better :

$$ab^{\frac{d}{2}} + b^{\frac{d}{2}} < b^{\frac{d}{2}}$$

Bi-directional
others
- Not optimal



AiML-
Notes-17-...

Informed Search:

- Uninformed: Goal not known, check if it is reached
 - Goal is known ahead time
 - heuristic func. is applied
- Best first
 A^*

[Lab using search algs]

16th Aug 2021 Evaluation

Group no: 22

Members: Shashwat.T (106119138), Hari Hara Sudhan (106119046)

- 1) Missing point: The one queue:
 - only for DFS and BFS it can be used
 - for UCS we have to only use priority queue
- 2) Apart from this the notes are perfect and covered every single line said in the lecture and PPT shown



Notes for
17-08-2021

FILE FOR EASIER VIEWING

Date: 17/08/2021

FILE FOR EASIER VIEWING

Date: 17/08/2021

Notes by : Group 25 Suseender (106119122) & Suraj (106119128)

Topics covered in class:

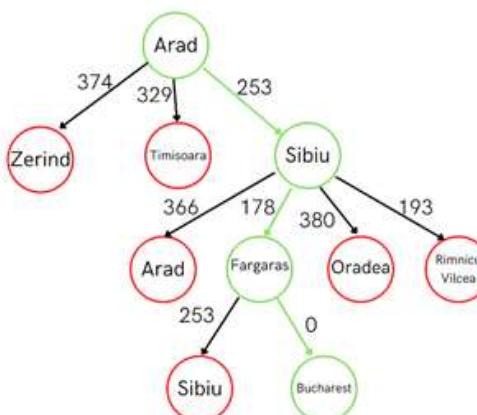
- Informed Search Strategies
 - Introduction to Heuristic functions
 - Greedy Best First Search
 - A* search algorithm

Informed Search Strategies:

- These are strategies that are based on the concept of Heuristic function:
 - Functions that decide on the efficiency of the algorithm where the goal is already known. It is a function that estimates how close a state is to a goal state. And designed for each problem individually.
 - The general approach is based on Best First Search (expand the best node first).
 - 2 types of approaches are the tree search or the graph search.
 - We go to the next node from the current node based on the evaluation function, f(n).
 - The Heuristic Function, h(n) is the cost of the lowest path from the start node 's' to the end goal 'g'.
 - One of the uses of these search strategies is in games where paths need to be found.
 - We use different ways of measuring distance such as Manhattan distance (*The Manhattan Distance between two points (X₁, Y₁) and (X₂, Y₂) is given by |X₁ - X₂| + |Y₁ - Y₂|*), Euclidean Distance for paths.
 - Heuristic function can be used to solve the Pancake Flipping problem.
 - Romania Map Example

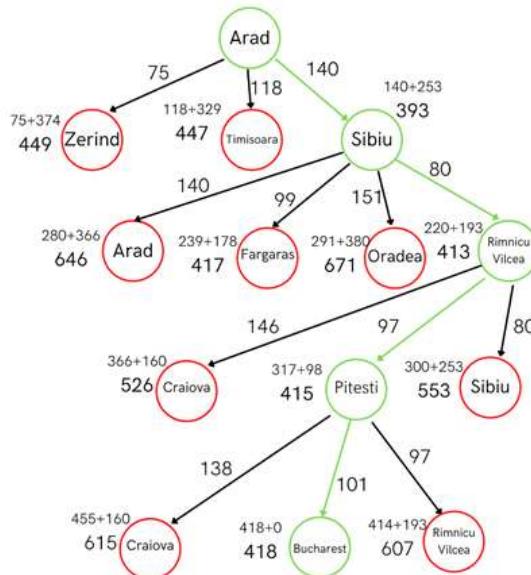


- For this example, let the start city be Arad and the end city be Bucharest. We have been given the distances between the cities and the straight-line distance between each city and the goal, Bucharest. (We can use haversine distance to estimate straight-line distance between 2 points using curvature. https://scikit-learn.org/stable/modules/generated/sklearn.metrics.pairwise.haversine_distances.html).
- **Using Greedy Best First Search**
 - Best First Search is based on the straight-line distance and the greedy approach uses a heuristic on the straight line distance to solve it.
 - In Greedy Best First Search you will look at the neighbouring nodes of the current node and then choose the one with the least straight-line distance to the goal city.
 - The heuristic is to estimate the distance to the nearest goal for each state.
 - Below is the Greedy Best First Search implementation for the Romanian map example:



- In the 1st step Sibiu is chosen as it has the least straight-line path to Bucharest and its nodes are expanded
- In the 2nd step Fagaras has the least straight-line path and it is expanded. This is done until Bucharest is reached.
- The final path is Arad->Sibiu->Fagaras->Bucharest, and the final cost comes to 450.

- In the 1st step Sibiu is chosen as it has the least straight-line path to Bucharest and its nodes are expanded.
 - In the 2nd step Fargaras has the least straight-line path and it is expanded. This is done until Bucharest is reached.
 - The final path is Arad->Sibiu->Fargaras->Bucharest, and the final cost comes to 450.
- Greedy is not very effective as we need straight line distance to all the nodes at the worst-case time complexity is (b^m) if all nodes are expanded.
- If goal is not properly defined, then it might lead to the wrong answer. The worst case is that it will become like a badly guided Depth First search.
- **A* Searching Algorithm**
 - Was used for a long time in games
 - Variations of A* are memory bounded A* (MA*) and simplified memory bounded A* (SMA*).
 - A* is similar to greedy but it is based on 2 functions rather than 1.
 - In greedy straight-line distance from current state to the goal state was considered but the path from source to the current state was not considered.
 - In A* we will consider 2 costs $g(n)$ and $h(n)$.
 - $g(n)$ is the cost to reach a particular node (path cost from start till current node).
 - $h(n)$ is the cost from the current node to the goal (this could be straight-line distance).
 - Therefore $f(n)$, the heuristic function can be stated as $f(n) = g(n) + h(n)$.
 - If UCS was a tortoise and greedy was a rabbit, then A* is both combined. UCS considers the cost of leaving a particular node and greedy considers the cost to reach a particular node. So, we can see how A* is a combination of these 2.
- The conditions for optimality for A* are:
 - Admissibility
 - The Admissible heuristic tries not to overestimate the cost to reach the goal.
 - A heuristic h is admissible if $0 \leq h(n) \leq h^*(n)$, where $h^*(n)$ is the true cost to a nearest goal.
 - Consistency
 - Also known as monotonicity.
 - A heuristic function $h(n)$ is called consistent if for all n and n' (*which is successor of n*) if the estimated cost for reaching the goal from n is not greater than the step cost of getting to n' + the cost of reaching the goal from n' .
 - $h(n) \leq c(n,a,n') + h(n')$ (*h(n) is cost to reach goal from n*)
 - It is a stronger condition.
 - It is trying to prove the correctness of you're A* algorithm.
 - It is used in graph search algorithms
- Below is the A* algorithm when used on the Romanian map example. At each node we will be calculating $g(n) + h(n)$.



- At each arrow the cost of going between the cities is written. The sum at each node is $g(n)+h(n)$ and the number below is the final sum $f(n)\{g(n)+h(n)\}$.
 - In the 1st step initial $g(n) = 0$ for starting city. For Zerind $g(n) = 75$ and $h(n) = 374$ and hence $f(n) = 75 + 374 = 449$. Similarly, for Timirosa and Sibiu $f(n)$ is calculated and least $f(n)$ from all the nodes is chosen.
 - We can see the final traversal route is Arad->Sibiu->Rimnicu Vilcea->Pitesti->Bucharest, and the final cost comes to 418.
 - We found that in the greedy method the final cost came to 450. Hence, we can see how the A* search algorithm can give a better result.
- When should A* terminate?
 - We normally stop when we come to a goal state
- Is A* Optimal?
 - A* works well if it is a tree but works less well in a graph where there maybe cycles in which it might get caught in travelling back and forth.
 - For A* Tree Search, it is optimal as it will reach the optimal fringe vertex first before a suboptimal one. Because all ancestors of optimal node will be reached before a suboptimal one and hence optimal solution will also be reached before the suboptimal solution.
- USC vs A*
 - While UCS expands in all directions equally, A* expands mainly toward the goal while hedging its bets to ensure optimality.
 - USC does not keep in mind the goal. A* moves such that it reaches the goal each time.
- Applications of A*

- While UCS expands in all directions equally, A* expands mainly toward the goal while hedging its bets to ensure optimality.
- USC does not keep in mind the goal. A* moves such that it reaches the goal each time.
- Applications of A*
 - A lot of video games
 - Paths and routes problems
 - Language analysis
 - Machine translation, etc
- 8 puzzle game can be solved using A* algorithm.

Evaluation for 17th August 2021 (Monday) Class

Done by Group 1 : 106119104-Rishi Sathish and 106119118-SreeGaneshTN

- Definition of Search Heuristics is missing .
- Proof of Optimality of A* Search is missing.
- Graph Example for Termination of A* is missing.
- UCS vs A* Contour is missing.
- Other the above mentioned points, everything is covered.

Class Notes by Group no: 33

Vishal Jaiswal : 106119144

Ritu Gain : 106119106

EVERWAY

Vishal Jaiswal , Ritu Gain
↓
106119144

↓
106119106 .

Informed search → Based upon heuristic function.

Search Heuristics:

Will decide on the efficiency of the algorithm where goal is reached.

General approach:

Best-first search (Which node will we expand first)
 → We will expand the best node first which is based upon the goal
 → Depends NOT ONLY upon the source.

In order to do that:

We can go for Tree search or Graph search } Choosing the approach is based upon:-

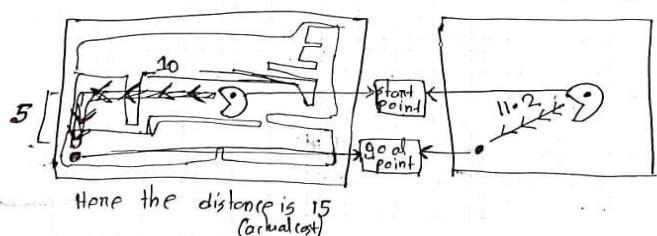
$h(n)$ = cost of lowest path from the start(s) Node 'n' to a goal state 'g'.

From the starting Node we'll use an evaluation function $f(n)$ which will be used for expanding the next Node and selecting the next Node.

A heuristic is

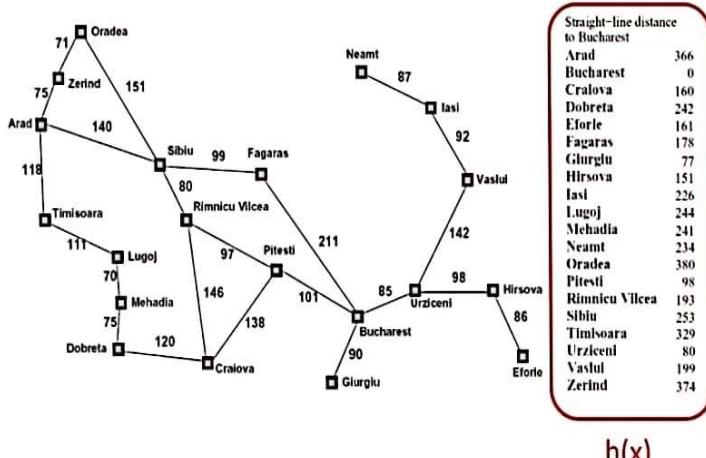
▪ A function that estimates how close a state is to a goal.
 ▪ designed for a particular search problem.

▪ Example: Manhattan distance, Euclidean distance for pathing (Routing).



11 to 2 is 5^{STD}
 STD (straight line distance)
 (Only if we're able to go)

Example: Heuristic Function



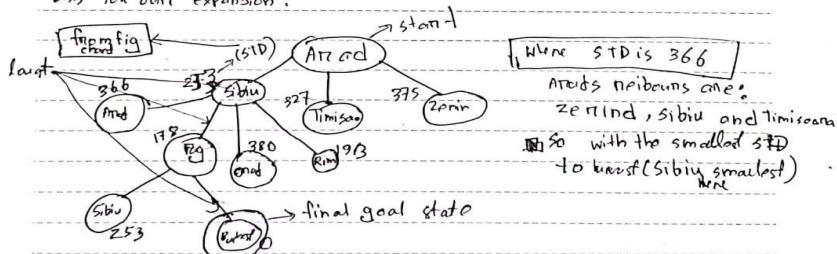
EVERWAY

DATE:

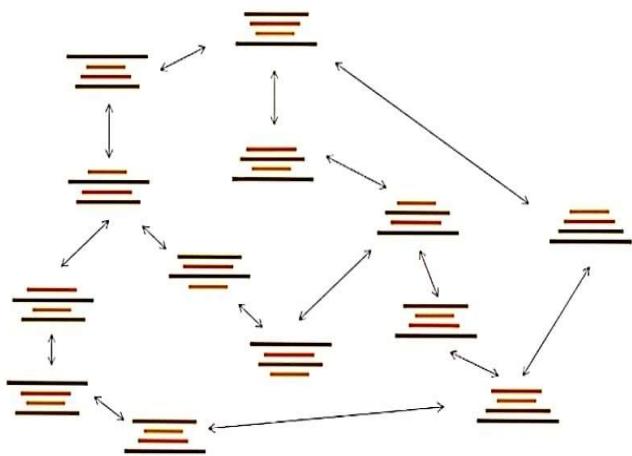
- So we have to get our ~~cost~~ Heuristic function such that the cost will going to be almost close to the SLD SLD.
- If we have landlord scenario we can use haversine distance in order to estimate the straightline distance between 2 points.
- Haversine distance is basically used to estimate the distance b/w 2 points in a geographical situation.
- We can use this function to solve Pancake problem.
→ We're going to flip the Pancakes such that the Pancakes are arranged from smaller to larger from top to bottom.

[Greedy Search / Greedy Best first search] (Not very effective) → os. we need a SLD for all nodes
→ based upon SLD as the heuristic function
[time complexity: $O(b^n)$]

- When we want to expand a particular Node the SLD is used as the basis for our expansion.



Example: Heuristic Function



EVERWAY

Greedy search

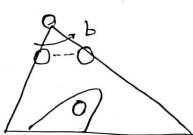
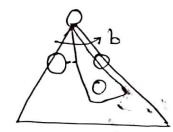
Strategy: expand a node that you think is closest to goal state.

* Heuristic: estimate of distance to nearest goal for each state

A common case:

* Best-first takes you straight to the (WRONG) goal. (sometimes)

Nonsensical case: like a badly-guided DFS



A* Search (greedy)

→ Memory bounded $A^*/(MA^*)$

→ Simplified " " $A^*/(SA^*)$

B based upon 2 functions as against a single heuristic function.

B The problem that discussed in Best-first search was that we always looked at the SLD from the new state to the goal state.

But the path cost from the source to the new state was not considered.

so we might have an alternate path which probably doesn't have good SLD But we might get less path cost. (if didn't consider that aspect of the alg.)

so, in A^* we'll consider 2 cost.

$g(n) \rightarrow$ cost to reach a particular node

$h(n) \rightarrow$ " " goal from the node to the goal (2nd distance cost)

so total function, $f(n) = g(n) + h(n)$

EVERWAY

DATE:

$g(n) \rightarrow$ Path cost from the start state to the Node n } c. const
 $h(n) \leftrightarrow$ SLD

$A^* \rightarrow UCS \rightarrow (g+h)$ (instead of g)

$UCS \rightarrow$ Imitate

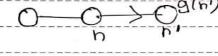
Breadth \rightarrow Rabbit

Optimality of A^*

Admissibility heuristic

\rightarrow tries to not overestimate the cost to reach goal.

$g(n) \rightarrow n, f(n) = g(n) + h(n)$



consistency \rightarrow stronger

in monotonicity

$A^* \rightarrow$ graph search algorithm.

$h(n)$ is consistent if $\forall n \& n'$ successor of n the estimated cost of reaching the goal from n isn't greater than the step cost of getting to n' + the cost of reaching the goal from n' .

$$h(n) \leq c(h, a, n') + h(n')$$

$h(n) \rightarrow$ cost to reach a goal state from n

$h(n') \rightarrow$ " " " " " " " "

- At every particular point of time when we're trying to go from one node to next node we will consider the path cost to reach the particular node + the cost to reach the goal state, so we have to cumulatively do that, to do that we should not over estimate or we shouldn't exceed exit the capacity - for example.

300 km 250 km

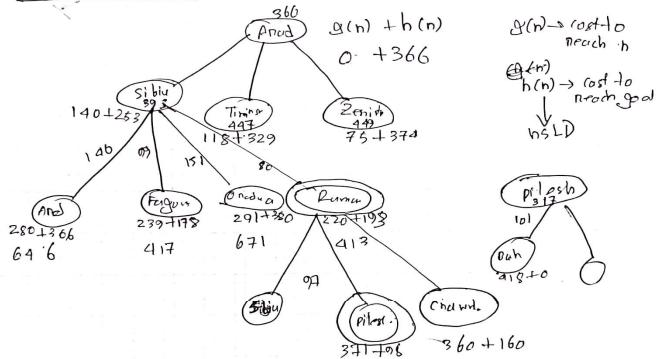


middle?

EVERWAY

Every consistent heuristic is called as consistent heuristic and admissible (no overestimate)

How A^* behaves:



$g(n) \rightarrow$ cost to reach n

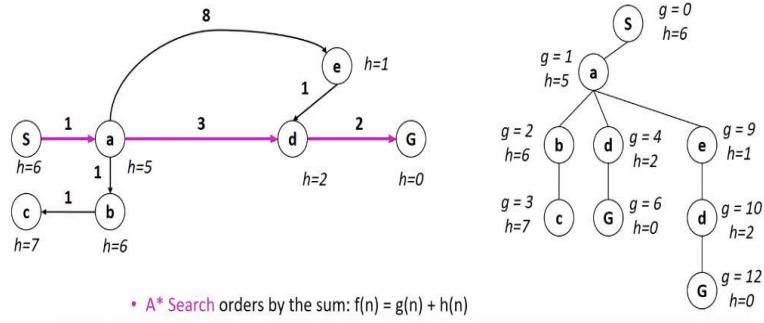
$h(n) \rightarrow$ cost to reach goal

↓

SLD

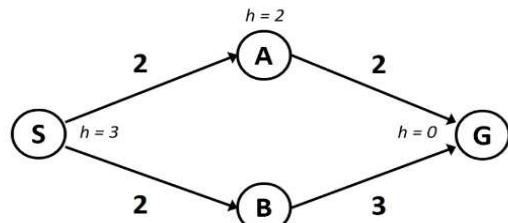
Combining UCS and Greedy

- Uniform-cost orders by path cost, or *backward cost* $g(n)$
- Greedy orders by goal proximity, or *forward cost* $h(n)$



When should A* terminate?

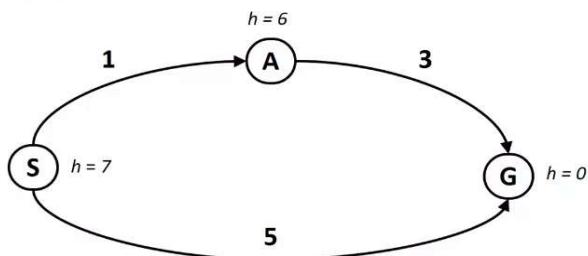
- Should we stop when we enqueue a goal?



- No: only stop when we ~~enqueue~~ a goal

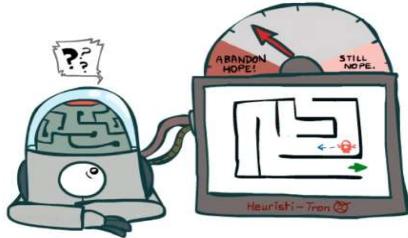
~~=====~~

Is A* Optimal?

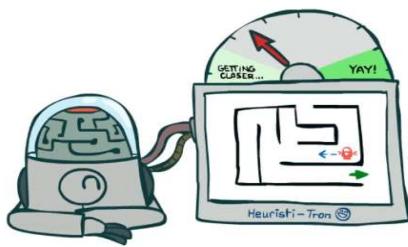


- What went wrong?
- Actual bad goal cost < estimated good goal cost
- We need estimates to be less than actual costs!

Idea: Admissibility



Inadmissible (pessimistic) heuristics break optimality by trapping good plans on the fringe



Admissible (optimistic) heuristics slow down bad plans but never outweigh true costs

Admissible Heuristics

- A heuristic h is *admissible* (optimistic) if:

$$0 \leq h(n) \leq h^*(n)$$

where $h^*(n)$ is the true cost to a nearest goal

$$h(n) \leq c(n, a, n') + h(n')$$

- Examples:



- Coming up with admissible heuristics is most of what's involved in using A* in practice.

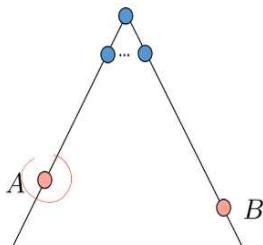
Optimality of A* Tree Search

Assume:

- A is an optimal goal node
- B is a suboptimal goal node
- h is admissible

Claim:

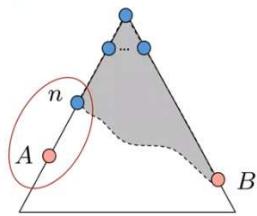
- A will exit the fringe before B



Optimality of A* Tree Search: Blocking

Proof:

- Imagine B is on the fringe
- Some ancestor n of A is on the fringe, too (maybe A!)
- Claim: n will be expanded before B
 1. $f(n)$ is less or equal to $f(A)$



$$f(n) = g(n) + h(n)$$

$$f(n) \leq g(A)$$

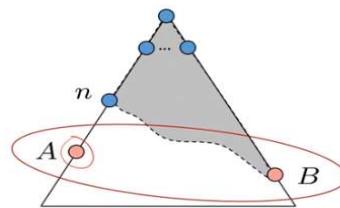
$$g(A) = f(A)$$

Definition of f-cost
Admissibility of h
 $h = 0$ at a goal

Optimality of A* Tree Search: Blocking

Proof:

- Imagine B is on the fringe
- Some ancestor n of A is on the fringe, too (maybe A!)
- Claim: n will be expanded before B
 1. $f(n)$ is less or equal to $f(A)$
 2. $f(A)$ is less than $f(B)$



$$g(A) < g(B)$$

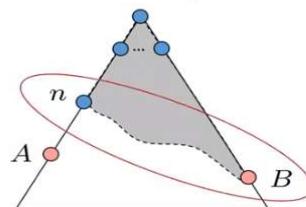
$$f(A) < f(B)$$

B is suboptimal
 $h = 0$ at a goal

Optimality of A* Tree Search: Blocking

Proof:

- Imagine B is on the fringe
- Some ancestor n of A is on the fringe, too (maybe A!)
- Claim: n will be expanded before B
 1. $f(n)$ is less or equal to $f(A)$
 2. $f(A)$ is less than $f(B)$
 3. n expands before B

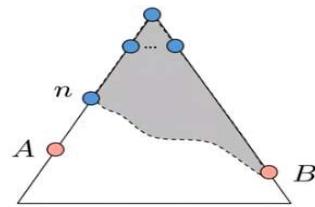


$$f(n) \leq f(A) < f(B)$$

Optimality of A* Tree Search: Blocking

Proof:

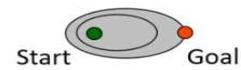
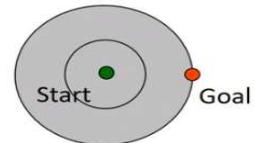
- Imagine B is on the fringe
- Some ancestor n of A is on the fringe, too (maybe A!)
- Claim: n will be expanded before B
 - 1. $f(n)$ is less or equal to $f(A)$
 - 2. $f(A)$ is less than $f(B)$
 - 3. n expands before B
- All ancestors of A expand before B
- A expands before B
- A* search is optimal



Properties of A*

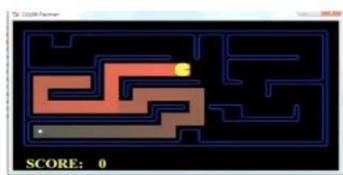
UCS vs A* Contours

- Uniform-cost expands equally in all “directions”
- A* expands mainly toward the goal, but does hedge its bets to ensure optimality

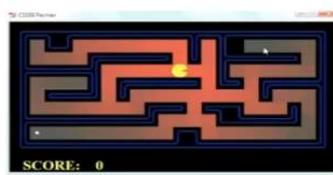


[Demo: contours UCS / greedy / A* empty (L3D1)]
[Demo: contours A* pacman small maze (L3D5)]

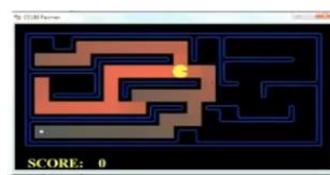
Comparison



Greedy



Uniform Cost



A*

A* Applications

- Video games
- Pathing / routing problems
- Resource planning problems
- Robot motion planning
- Language analysis
- Machine translation
- Speech recognition
- ...

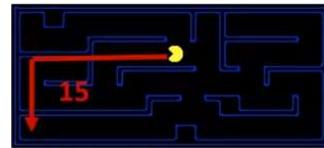


[Demo: UCS / A* pacman tiny maze (L3D6,L3D7)]

[Demo: guess algorithm Empty Shallow/Deep (L3D8)]

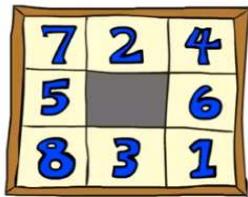
Creating Admissible Heuristics

- Most of the work in solving hard search problems optimally is in coming up with admissible heuristics
- Often, admissible heuristics are solutions to *relaxed problems*, where new actions are available

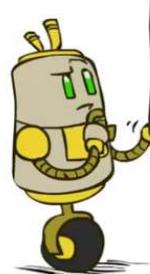


- Inadmissible heuristics are often useful too

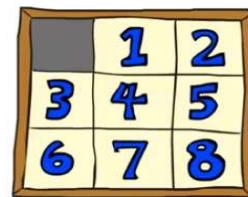
Example: 8 Puzzle



Start State



Actions

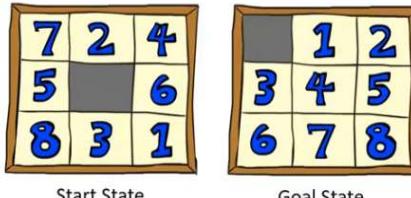
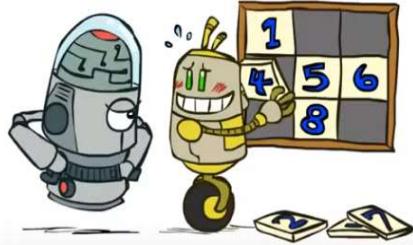


Goal State

- What are the states?
- How many states?
- What are the actions?
- How many successors from the start state?
- What should the costs be?

8 Puzzle I

- Heuristic: Number of tiles misplaced
- Why is it admissible?
- $h(\text{start}) = 8$
- This is a relaxed-problem heuristic



Average nodes expanded when the optimal path has...		
	...4 steps	...8 steps
UCS	112	6,300
TILES	13	39
		3.6×10^6
	...12 steps	227

8 Puzzle II

- What if we had an easier 8-puzzle where any tile could slide any direction at any time, ignoring other tiles?

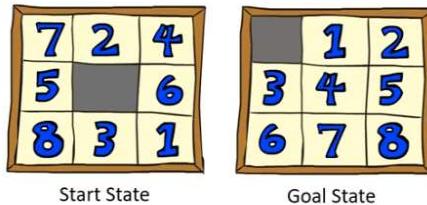
- Total Manhattan distance

City block distance

- Why is it admissible?

$$3 + 1 + 2 + \dots = 18$$

- $h(\text{start}) =$



Average nodes expanded when the optimal path has...		
	...4 steps	...8 steps
TILES	13	39
MANHATTAN	12	25
		73

8 Puzzle III

- How about using the *actual cost* as a heuristic?

- Would it be admissible?
- Would we save on nodes expanded?
- What's wrong with it?



- With A*: a trade-off between quality of estimate and work per node

- As heuristics get closer to the true cost, you will expand fewer nodes but usually do more work per node to compute the heuristic itself



AI_ML_NOT

ES

PDF File for easy viewing

23/8/21

will it always lead to a complete solution?
↳ Yes, but the problem is, memory, if the destination is too far.

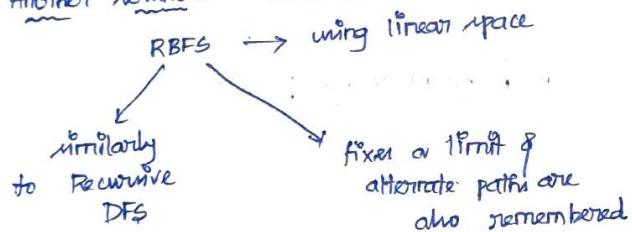
so, the solution is:

→ memory bounded heuristic search
It is similar to iterative deepening. (IDA*)

$$h(n) = f(n) + g(n), \text{ [with depth fixed]}$$

→ thus the amount of memory req to store intermediate states is reduced.

Another solution: Recursive Best first search



pseudocode:

```
return RBFS (problem, make-node (prob, initialstate α))
```

RBFS (—, —, —)

if problem goal test (node.state) return s/n;
succsor [] // enough memory

for each action in problem.action (node.state)

f-limit

initialstate
α

if (succsor is empty) return false

the (Π is successor & available)

& s in memory

$$s.f = \max(s.g + s.h, \text{node}.f)$$

do

best = lowest f .value node in successor

if $\text{best}.f > f\text{-limit}$ return false, $\text{best}.f$;

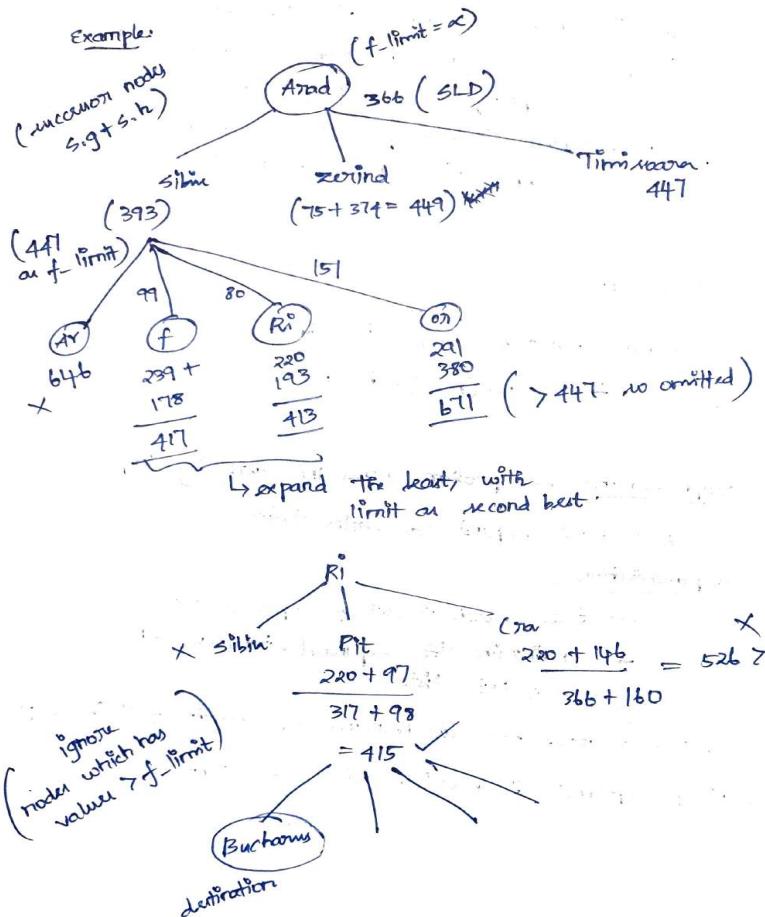
alternative node = record lowest f .value of
successor

result, best. f = RBFS(problem, best,
 $\min(f\text{-limit},$
 $\text{alternative node}))$

if result \neq failure

return result as solution.

successors can be found using priority-queue.



simplified memory bounded A^* (smA*)

↳ new node can't be added without dropping an existing node

→ backtracking is not possible as we have dropped few nodes from trees.

Trivial Heuristics

dominance function: $h_a \geq h_c$ if

$$\nexists : h_a(n) > h_c(n)$$

graph search: → problem comes in detecting loops

↳ no never expand a vertex twice

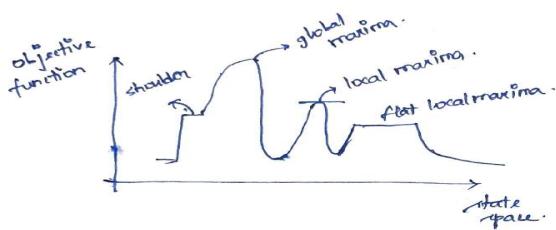
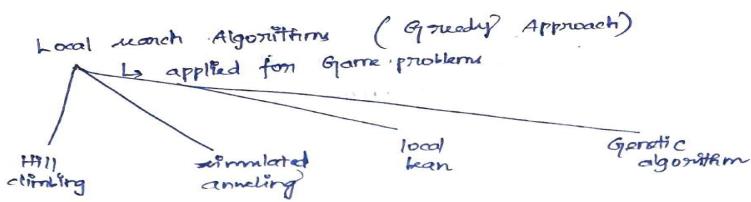
Implementation:

→ tree search + set of expanded states

→ maintain the expanded states on SET
as not lists

A^* → looks for loc minima (no might fail in graphs)

→ A^* is optimal if heuristic is consistent



Hill-climbing search:

→ climb - uphill

```

funct hc (problem)
{
  current ← make-node (prob, initial-state)
  do
    neighbours ← a high valued members of
    current
    if neighbours.value ≤ current.value return
      current.state.
    current ← neighbours.
}
heuristic setting : local max : greater than current
but not Global maximum
ridge : A result of all local maxima
in close to each other.
  
```

plateau → flat area of the state space
 → flat - local maxima / shoulder
 → no hill climb

these algorithms